

Solving Grid-World Problem with Reinforcement Learning

Jean Lucien Randrianantenaina
Applied Mathematics, Machine Learning and Artificial Intelligence
Stellenbosch University

Abstract—

I. INTRODUCTION

II. PROBLEM STATEMENT AND MODELLING

We aim to build an agent that solves an empty 6×6 grid world, e.g., reaching the green cell as illustrated in Figure 1. The agent (red triangle) interacts with the environment through observations, actions and rewards. We use the `mini-grid` API from the Farama foundation to simulate the environment.




Fig. 1: The MiniGrid-Empty-8x8-v0 environment

Depending on our configuration, the state can be a tensor or an image which is returned after each action. The agent has eight possible actions, but we only use three: turn left and right and move forward. In this setting, we deal with an episodic task, and for each episode, the agent can perform at most 256 steps.

If the agent reaches the target, it receives a reward r defined by:

$$r = 1 - 0.9 \times \frac{\text{Number of steps}}{\text{Maximum number of steps}} \quad (1)$$

otherwise, the reward is zero.

Therefore our task is to make the agent learn how to move in this grid in such a way that it maximizes its reward. For




Fig. 2: Reinforcement learning Model

this particular environment, the agent can get at most a reward of 0.9613 with 11 steps. But for this work, we will consider 0.9578 as the optimal reward corresponding to 12 steps.

III. REINFORCEMENT LEARNING TECHNIQUE

In this era, reinforcement learning is widely used to solve the problem of programming intelligent agents that learn how to optimise a specific task such as playing a game, or autonomous vehicle. This is done by learning a policy for sequential decision problems that maximize the discounted cumulative future reward:

$$G_t \triangleq \sum_{k=t+1}^T \gamma^{k-t-1} R_k \quad (2)$$

Where R is the reward, $\gamma \in [0, 1]$ is the discount factor that controls the importance from immediate reward (near to 0) or future reward (near to 1). The agent learns the optimal policy without being explicitly told if its actions are good or bad using the reinforcement learning model depicted in Figure 2.

Several algorithms/techniques can be used to solve our problem with reinforcement learning, in the next section we will look:

- Q-Learning
- SARSA (State-Action-Reward-State-Action)
- Deep Q-Network
- Deep Q-Network with RGB Image techniques.

The algorithm 1 show the skeleton shared by these RL algorithms. In addition to that, we also use the ϵ -greedy strategy to select the action to be performed by the agent. The agent uniformly samples an action from the possible action for the given state with a probability ϵ and uses the optimal action

Parameters: . . .

```

foreach episode do
  (Re)Initialize the environment
  foreach step do
    Act and consider the observation and reward
    Steps specific to each methods
    if done or truncated then
      | Some steps for metric and monitoring
    end
  end
end

```

Algorithm 1: RL Algorithm Skeleton

for the given state with a probability $1 - \epsilon$. In the first case, we say that the agent is exploring, and in the second case we say that the agent is exploiting.

For this work we will start with a higher value of ϵ to allow the agent to explore the environment, then we decrease it slowly until we reach a small value, this will be controlled by :

$$\epsilon_t = \epsilon_{min} + (\epsilon_{max} - \epsilon_{min}) \exp \left\{ -\frac{t}{\Delta} \right\} \quad (3)$$

Where Δ is the decay rate, a higher value corresponds to a slow decrease, and a small value corresponds to a fast decrease.

IV. TABULAR METHODS: Q-LEARNING AND SARSA

The tabular method refers to the creation of a table called Q-table, containing the value of $Q(S_t = s, A_t = a)$. This value indicates how good is acting a on the state s .

For the two algorithms that we present in this section, the observation (a $7 \times 3 \times 3$ array) will be converted into a string to create an MD5 hash to represent the given state.

A. SARSA

The SARSA algorithm is an on-policy method, which means it iteratively refines a single policy, that also generates control action within the environments. The update rule is defined by:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)] \quad (4)$$

From the equation (4), we obtain the Algorithm 2 for the SARSA learning.

B. Q-Learning

In contrast to SARSA, the Q-Learning algorithm is an off-policy algorithm, which means we optimise the target policy using different policies. The update rule is given by

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)] \quad (5)$$

Since we have the Algorithm 3 for Q-learning methods.

These two methods are powerful in dealing with a few states and actions, but not very efficient for numerous states and actions. We can replace the table with a function approximator as we will see in the next section.

Parameters: Step size $\alpha \in]0, 1]$, small $\epsilon > 0$

Initialize $Q(s, a)$ for all $s \in \mathcal{S}^+$ and $a \in \mathcal{A}(s)$ arbitrarily, except that $Q(\text{terminal} - \text{state}, \cdot) = 0$

```

foreach episode do
  Initialize  $S$ 
  Choose  $A$  from  $S$  using policy derived from  $Q$  ( $\epsilon$ -greedy)
  foreach step until S is a terminal state do
    Take the action  $A$ , observe  $R, S'$ 
    Choose  $A'$  from  $S'$  using policy derived from  $Q$  ( $\epsilon$ -greedy)
     $Q(S, A) \leftarrow Q(S, A) + \alpha [R_{t+1} + \gamma Q(S', A') - Q(S_t, A_t)]$ 
     $S \leftarrow S'$ 
     $A \leftarrow A'$ 
  end
end

```

Algorithm 2: SARSA: on-policy learners to estimate the optimal Q-table

Parameters: Step size $\alpha \in (0, 1]$, small $\epsilon > 0$

Initialize $Q(s, a)$ for all $s \in \mathcal{S}^+$ and $a \in \mathcal{A}(s)$ arbitrarily, except that $Q(\text{terminal} - \text{state}, \cdot) = 0$

```

foreach episode do
  Initialize  $S$ 
  Choose  $A$  from  $S$  using  $\epsilon$ -greedy
  foreach step until S is a terminal state do
    Take the action  $A$ , observe  $R, S'$ 
    Choose  $A'$  from  $S'$  using  $\epsilon$ -greedy
     $Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$ 
     $S \leftarrow S'$ 
  end
end

```

Algorithm 3: Q-learning algorithm to estimate the optimal Q-table

V. FUNCTION APPROXIMATION: DEEP Q-NETWORK

The idea of function approximation is to replace the value function with its approximation instead of using a look-up table. The function can take

- the state and action as input, and $q(s, a)$ as output,
- or the state as input and $q(s, a)$ as output.

This is illustrated in the figure 3.




Fig. 3: Illustration of function approximation

In our case, we will use a neural network as a function approximation. Thus we will consider two types.

A feed-forward neural network (MLP) which takes a vector $\mathbf{u} \in \mathbb{R}^{49 \times 1}$ (state) as input and $q(s, a_1), q(s, a_2), q(s, a_3)$ as output.

Similarly, Convolutional Neural networks (CNN), take a stack of four frames (images) at four successive time steps, e.g. a $4 \times 56 \times 56$ tensor as input and the same output as the MLP. The architecture of these neural networks are described in respectively in REF-MLP and REF-CNN.

As we do not have any labels to train the networks, we use two networks: the `policy_net` and `target_net`. The first one is optimised with the Adam optimizer, while the second one is fixed, and used to generate a sort of label for the `policy_net`. We also synchronize the weight of the two networks in a regular period of the training, to move toward the optimal values.

To train these networks, we need a deque (double end-queue) that stores a given number of the experience of the agent formed by the current state, action, next state, and reward, we call this a replay memory \mathcal{D} . As new experiences are added the oldest data is pushed out of the deque. Once we reach have batch size (enough number) of experience, we start to sample from \mathcal{D} and optimize the `policy_net` parameter by minimizing the Mean Squared Error (MSE) given by:

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N \left[R_i + \gamma \max_{a'} \hat{Q}(s'_i, a', \mathbf{w}^-) - \hat{Q}(s_i, a, \mathbf{w}) \right]^2 \quad (6)$$

where $R_i + \gamma \max_{a'} \hat{Q}(s'_i, a', \mathbf{w}^-)$ is computed using `target_net`, and $\hat{Q}(s'_i, a', \mathbf{w}^-)$ is 0 if s'_i is terminal state, while $\hat{Q}(s_i, a, \mathbf{w})$ is computed with `policy_net`.

Note that, the input of these networks (observation) is normalized by rescaling their value between $[-1, 1]$. In particular, we convert the RGB image into a grey-scale image for the CNN architecture. We can summarize the whole process in the algorithm 4.

VI. HYPER-PARAMETERS TUNING AND EVALUATION

A. Hyper-parameters

like all machine learning problems we have several parameters to tune to get the optimal results. Grid search is one of the best methods, it iterates through all the possible combinations of predefined parameters set. The drawback of this method is the time complexity which can explode considerably when we hyper-parameter space is big and the algorithm is quite slow.

So, instead of using that approach, we will use trial and error to find good hyper-parameters such as γ , α , and the number of episodes. To do that we start we some values, and we adjust these parameters according to the obtained results.

B. Model evaluation

To evaluate the model, we run the agents in the environment for 1000 episodes, then compute the completion rate, the reward average, and the averaged steps for each method.

We also investigate these values for the training process to assess the speed and efficiency of each method, in addition to the plot of the loss and accumulated rewards.

```

Initialize policy_net and target_net
Initialize the environments
Set the decay rate for the epsilon decreasing
Set the updating period of target_net
Set the total step to 0
Create a replay memory D
foreach episode do
    Set step to 0
    Make a new episode
    Observe the first state
    while not( done or truncated) do
        Choose A from S using policy derived from Q
        (ε-greedy)
        Increment the total training step
        Execute A, observe R and the new state S'
        Store Transition < S, A, S', R > in D
        Compute the L and do a gradient descent step
        if updating period then
            Copy the policy_net parameter to
            target_net
        end
    end
end

```

Algorithm 4: Training a DQN to estimate the optimal policy

VII. RESULTS AND DISCUSSION

In this section we discuss about our results and present them. The smoothed plot were smoothed using the exponential moving average with a factor of 0.7.

A. Q-Learning




Fig. 4: Q-Learning per episode training metrics

B. SARSA-Learning




Fig. 5: SARSA-Learning per episode training metrics

C. Deep Q-Network Learning




Fig. 6: DQN-Learning per episode training metrics

D. Deep Q-Network Learning with RGB Image





Fig. 7: DQN-Learning with RGB Image technique per episode training metrics



VIII. CONCLUSION