

Results

July 31, 2024

```
[1]: import json
import pickle

import matplotlib.animation as animation
import matplotlib.pyplot as plt
from IPython.display import HTML


from Utility.Plots import plot_live_data, plot_train_data, plot_steps_taken
from Utility.TabularEpsilonStrategy import EpsilonGreedy

%matplotlib inline
%config InlineBackend.figure_format = 'retina'
```

0.1 Q-Learning

```
[2]: with open('Q_Learning/QLearning_live_plot.json', 'r') as f:
    data = json.load(f)

plt.figure(figsize=(16, 5))
plt.subplot(131)
plot_live_data(data['episode'], data['rewards'], epsilon=data['epsilons'],
               title={'title': "Q-Learning Performance", 'raw': 'Raw Rewards',
                       'smooth': 'Smoothed Rewards',
                       'acc' : 'Average Accumulated Rewards', 'xlab': 'Episode',
                       'ylab': 'Reward'})
plt.subplot(132)
plot_train_data(data['episode'], data['losses'],
               {'title': "Q-Learning Training", 'raw': 'Raw TD-Error**2',
                 'smooth': 'Smoothed TD-Error**2',
                 'acc' : 'Average Accumulated TD-Error**2', 'xlab': 'Episode',
                 'ylab': 'Loss'})
plt.subplot(133)
plot_steps_taken(data['episode'], data['steps_taken'], "Q-Learning")
plt.tight_layout()
plt.savefig('report/figures/QLearning_episode.pdf')
plt.savefig('report/figures/QLearning_episode.png')
plt.show()
```



```
[3]: with open('Q_Learning/QLearning_train.json', 'r') as f:
    data = json.load(f)


plt.figure(figsize=(16, 5))

plt.subplot(121)
plot_live_data(data['steps_done'], data['rewards'], avg=False,
               title={'title': "Rewards per Step", 'raw': 'Raw Rewards',
                           'smooth': 'Smoothed Rewards',
                           'acc' : 'Cumulative Average Rewards', 'xlab': 'Episode',
                           'ylab': 'Rewards'})

plt.subplot(122)
plot_train_data(data['steps_done'], data['td_error_sq'], avg=False,
                title={'title' : "TD-Error**2 per Step", 'raw': 'Raw_
                           TD-Error**2',
                           'smooth': 'Smoothed TD-Error**2', 'acc': 'Cumulative_
                           TD-Average TD-Error**2',
                           'xlab' : 'Episode', 'ylab': 'TD-Error**2'})

plt.tight_layout()
plt.savefig('report/figures/QLearning_steps.pdf')
plt.savefig('report/figures/QLearning_steps.png')

plt.show()
```



```
[4]: rewards = data['rewards']
steps_taken = data['steps_taken']
finish_counter = data['finish_counter']
episodes = len(rewards)
print('===== TRAIN SUMMARY =====')
print(f'Episodes      : {episodes}')
print(f"Completion rate: {finish_counter / episodes}")
print(f"Average Reward  : {sum(rewards) / episodes:.3f}")
print(f"Average steps   : {sum(steps_taken) / episodes:.3f}")
```

```
===== TRAIN SUMMARY =====
Episodes      : 512
Completion rate: 0.9296875
Average Reward : 0.808
Average steps   : 52.510
```

```
[5]: from Q_Learning.QLearning import eval as qlearning_eval
from Utility.MiniGrid import MiniGridHash

env = MiniGridHash(render_mode='rgb_array')
with open('Q_Learning/q_learning_table.pkl', 'rb') as f:
    qlearning_table = pickle.load(f)
qlearning_eval(env, qlearning_table, EpsilonGreedy(0, 3), 1000)
```

```
pygame 2.6.0 (SDL 2.28.4, Python 3.10.14)
Hello from the pygame community. https://www.pygame.org/contribute.html
Evaluation...
===== EVALUATION SUMMARY =====
Evaluation episodes: 1000
Completion rate      : 1.0
Average Reward       : 0.958
Average steps        : 12.000
```

```
[6]: imgs = []
fig = plt.figure(figsize=(5, 5))
strategy = EpsilonGreedy(0, 3)
current_state = env.reset()
plt.axis('off')
im = plt.imshow(env.render(), animated=True)
imgs.append([im])
for step in range(env.maxSteps):
    action = strategy.select_action(current_state, qlearning_table)
    next_state, reward, done, truncated = env.step(action)
    im = plt.imshow(env.render(), animated=True)
```

```

    imgs.append([im])
    if done or truncated:
        break
    current_state = next_state
ani = animation.ArtistAnimation(fig, imgs, interval=100, repeat_delay=1000)
ani.save('Q_Learning/q_learning.gif', writer='imagemagick')
plt.close()
HTML(ani.to_jshtml())

```

MovieWriter imagemagick unavailable; using Pillow instead.

[6]: <IPython.core.display.HTML object>


0.2 SARSA

```

[7]: with open('SARSA/SARSA_live_plot.json', 'r') as f:
    data = json.load(f)

plt.figure(figsize=(16, 5))
plt.subplot(131)
plot_live_data(data['episode'], data['rewards'], epsilon=data['epsilons'],
               title={'title': "SARSA-Learning Performance", 'raw': 'Raw Rewards',
                       'smooth': 'Smoothed Rewards',
                       'acc' : 'Average Accumulated Rewards', 'xlab': 'Episode',
                       'ylab': 'Reward'})
plt.subplot(132)
plot_train_data(data['episode'], data['losses'],
                {'title': "SARSA-Learning Training", 'raw': 'Raw TD-Error**2',
                  'smooth': 'Smoothed TD-Error**2',
                  'acc' : 'Average Accumulated TD-Error**2', 'xlab': 'Episode',
                  'ylab': 'Loss'})
plt.subplot(133)
plot_steps_taken(data['episode'], data['steps_taken'], "SARSA-Learning")
plt.tight_layout()
plt.savefig('report/figures/SARSA_Learning_episode.pdf')
plt.savefig('report/figures/SARSA_Learning_episode.png')
plt.show()

```




```
[8]: with open('SARSA/SARSA_train.json', 'r') as f:
      data = json.load(f)

plt.figure(figsize=(16, 5))

plt.subplot(121)
plot_live_data(data['steps_done'], data['rewards'], avg=False,
               title={'title': "Rewards per Step", 'raw': 'Raw Rewards',
               ↪ 'smooth': 'Smoothed Rewards',
               'acc' : 'Cumulative Average Rewards', 'xlab': 'Episode',
               ↪ 'ylab': 'Rewards'})

plt.subplot(122)
plot_train_data(data['steps_done'], data['td_error_sq'], avg=False,
                title={'title' : "TD-Error**2 per Step", 'raw': 'Raw
                ↪ TD-Error**2',
                'smooth': 'Smoothed TD-Error**2', 'acc': 'Cumulative
                ↪ Average TD-Error**2',
                'xlab' : 'Episode', 'ylab': 'TD-Error**2'})

plt.tight_layout()
plt.savefig('report/figures/SARSA Learning_steps.pdf')
plt.savefig('report/figures/SARSA Learning_steps.png')
plt.show()
```



```
[9]: rewards = data['rewards']
steps_taken = data['steps_taken']
finish_counter = data['finish_counter']
episodes = len(rewards)
print('===== TRAIN SUMMARY =====')
print(f'Episodes      : {episodes}')
print(f'Completion rate: {finish_counter / episodes}')
```

```
print(f"Average Reward : {sum(rewards) / episodes:.3f}")
print(f"Average steps : {sum(steps_taken) / episodes:.3f}")
```

```
===== TRAIN SUMMARY =====
Episodes      : 600
Completion rate: 0.9566666666666667
Average Reward : 0.844
Average steps  : 43.185
```

0.2.1 Evaluation

```
[10]: from SARSA.SARSA import eval as sarsa_eval
      from Utility.MiniGrid import MiniGridHash

      env = MiniGridHash(render_mode='rgb_array')
      with open('SARSA/sarsa_learning_table.pkl', 'rb') as f:
          sarsa_table = pickle.load(f)
      sarsa_eval(env, qlearning_table, EpsilonGreedy(0, 3), 1000)
```

```
Evaluation...
===== EVALUATION SUMMARY =====
Evaluation episodes: 1000
Completion rate      : 1.0
Average Reward       : 0.958
Average steps        : 12.000
```

```
[11]: imgs = []
      fig = plt.figure(figsize=(5, 5))
      strategy = EpsilonGreedy(0, 3)
      current_state = env.reset()
      plt.axis('off')
      im = plt.imshow(env.render(), animated=True)
      imgs.append([im])
      for step in range(env.maxSteps):
          action = strategy.select_action(current_state, sarsa_table)
          next_state, reward, done, truncated = env.step(action)
          im = plt.imshow(env.render(), animated=True)
          imgs.append([im])
          if done or truncated:
              break
          current_state = next_state
      ani = animation.ArtistAnimation(fig, imgs, interval=100, repeat_delay=1000)
      ani.save('SARSA/sarsa.gif', writer='imagemagick')
      plt.close()
      HTML(ani.to_jshtml())
```


MovieWriter imagemagick unavailable; using Pillow instead.

```
[11]: <IPython.core.display.HTML object>
```

0.3 DQN

```
[12]: with open('DQN/DQN_live_plot.json', 'r') as f:
        data = json.load(f)

plt.figure(figsize=(16, 5))
plt.subplot(131)
plot_live_data(data['episode'], data['rewards'], epsilon=data['epsilons'],
               title={'title': "DQN-Learning Performance", 'raw': 'Raw Rewards',
                       'smooth': 'Smoothed Rewards',
                       'acc' : 'Average Accumulated Rewards', 'xlab': 'Episode',
                       'ylab': 'Reward'})
plt.subplot(132)
plot_train_data(data['episode'], data['losses'],
               {'title': "DQN-Learning Training", 'raw': 'Raw TD-Error**2',
               'smooth': 'Smoothed TD-Error**2',
               'acc' : 'Average Accumulated TD-Error**2', 'xlab': 'Episode',
               'ylab': 'Loss'})
plt.subplot(133)
plot_steps_taken(data['episode'], data['steps_taken'], "DQN-Learning")
plt.tight_layout()
plt.savefig('report/figures/DQNLearning_episode.pdf')
plt.savefig('report/figures/DQNLearning_episode.png')
plt.show()
```



```
[13]: with open('DQN/DQN_train.json', 'r') as f:
        data = json.load(f)

plt.figure(figsize=(16, 5))

plt.subplot(121)
plot_live_data(data['steps_done'], data['rewards'], avg=False,
```

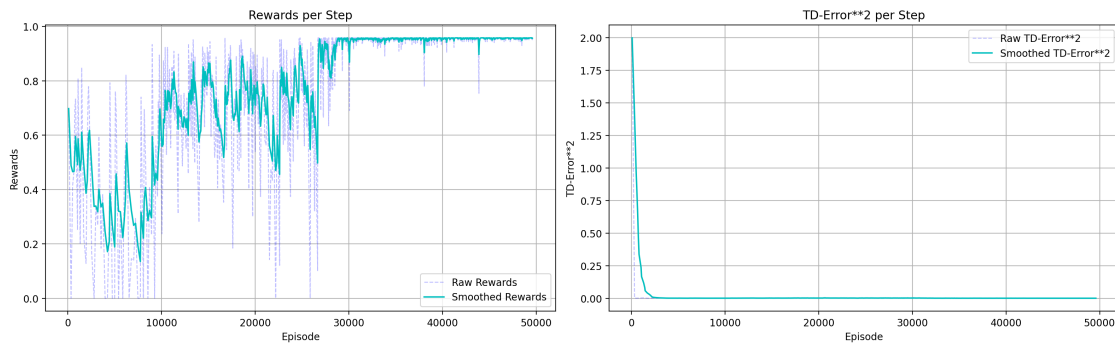
```

        title={'title': "Rewards per Step", 'raw': 'Raw Rewards',
        ↪ 'smooth': 'Smoothed Rewards',
        'acc' : 'Cumulative Average Rewards', 'xlab': 'Episode',
        ↪ 'ylab': 'Rewards'})

plt.subplot(122)
plot_train_data(data['steps_done'], data['td_error_sq'], avg=False,
        title={'title' : "TD-Error**2 per Step", 'raw': 'Raw
        ↪ TD-Error**2',
        'smooth': 'Smoothed TD-Error**2', 'acc': 'Cumulative
        ↪ Average TD-Error**2',
        'xlab' : 'Episode', 'ylab': 'TD-Error**2'})

plt.tight_layout()
plt.savefig('report/figures/DQNLearning_steps.pdf')
plt.savefig('report/figures/DQNLearning_steps.png')
plt.show()

```



```

[14]: rewards = data['rewards']
steps_taken = data['steps_taken']
finish_counter = data['finish_counter']
episodes = len(rewards)
print('===== TRAIN SUMMARY =====')
print(f'Episodes      : {episodes}')
print(f"Completion rate: {finish_counter / episodes}")
print(f"Average Reward : {sum(rewards) / episodes:.3f}")
print(f"Average steps  : {sum(steps_taken) / episodes:.3f}")

```

```

===== TRAIN SUMMARY =====
Episodes      : 2000
Completion rate: 0.992
Average Reward : 0.912
Average steps  : 24.794

```


0.4 Evaluation

```
[15]: from DQN.DQN import eval as dqn_eval
      from DQN.DQN import DQN
      from Utility.DQNEpsilonStrategy import EpsilonGreedy as DQNEpsilonGreedy
      from Utility.Minigrid import MiniGridRaw
      import torch
      from torchinfo import summary

      env_dqn = MiniGridRaw(render_mode='rgb_array')
      policy_net = DQN(env_dqn.numStates, env_dqn.numActions, (64, 32))
```

```
[16]: a = env_dqn.reset()
      print(summary(policy_net, input_data=a, col_names=['input_size', 'output_size', 'num_params', 'trainable']))
```

```
=====
=====
Layer (type:depth-idx)      Input Shape      Output Shape
Param #                    Trainable
=====
=====
DQN                        [1, 49]          [1, 3]
--                          True
  Sequential: 1-1          [1, 49]          [1, 3]
--                          True
    Linear: 2-1            [1, 49]          [1, 64]
3,200                      True
    ReLU: 2-2              [1, 64]          [1, 64]
--                          --
    Linear: 2-3            [1, 64]          [1, 32]
2,080                      True
    ReLU: 2-4              [1, 32]          [1, 32]
--                          --
    Linear: 2-5            [1, 32]          [1, 3]
99                          True
=====
=====
Total params: 5,379
Trainable params: 5,379
Non-trainable params: 0
Total mult-adds (M): 0.01
=====
=====
Input size (MB): 0.00
Forward/backward pass size (MB): 0.00
Params size (MB): 0.02
Estimated Total Size (MB): 0.02
```

```
=====
=====
```

```
[17]: policy_net.load_state_dict(torch.load('DQN/dqn.pth'))
      dqn_eval(env_dqn, policy_net, DQNEpsilonGreedy(0, 3), 1000)
```

Evaluation...

===== EVALUATION SUMMARY =====

Evaluation episodes: 1000

Completion rate : 1.0

Average Reward : 0.958

Average steps : 12.000

```
[18]: imgs = []
      fig = plt.figure(figsize=(5, 5))
      strategy_dqn = DQNEpsilonGreedy(0, 3)
      current_state = env_dqn.reset()
      plt.axis('off')
      imgs.append([plt.imshow(env_dqn.render(), animated=True)])
      for step in range(env.maxSteps):
          action = strategy_dqn.select_action(current_state, policy_net)
          next_state, reward, done, truncated = env_dqn.step(action)
          imgs.append([plt.imshow(env_dqn.render(), animated=True)])
          if done or truncated:
              break
          current_state = next_state
      ani = animation.ArtistAnimation(fig, imgs, interval=100, repeat_delay=1000)
      ani.save('DQN/dqn.gif', writer='imagemagick')
      plt.close()
      HTML(ani.to_jshtml())
```

MovieWriter imagemagick unavailable; using Pillow instead.

```
[18]: <IPython.core.display.HTML object>
```

0.5 DQN RGB Technique


```
[19]: with open('DQN/DQNIMAGE_live_plot.json', 'r') as f:
      data = json.load(f)

      plt.figure(figsize=(16, 5))
      plt.subplot(131)
      plot_live_data(data['episode'], data['rewards'], epsilon=data['epsilons'],
                     title={'title': "DQNRGB-Learning Performance", 'raw': 'Raw',
                             'Rewards': 'Smoothed Rewards',
                             'acc' : 'Average Accumulated Rewards', 'xlab': 'Episode', 'ylab': 'Reward'})
      plt.subplot(132)
```

```

plot_train_data(data['episode'], data['losses'],
                {'title': "DQNRGB-Learning Training", 'raw': 'Raw TD-Error**2',
                 'smooth': 'Smoothed TD-Error**2',
                 'acc' : 'Average Accumulated TD-Error**2', 'xlab': 'Episode',
                 'ylab': 'Loss'})
plt.subplot(133)
plot_steps_taken(data['episode'], data['steps_taken'], "DQNRGB-Learning")
plt.tight_layout()
plt.savefig('report/figures/DQNRGBLearning_episode.pdf')
plt.savefig('report/figures/DQNRGBLearning_episode.png')
plt.show()

```



```

[20]: with open('DQN/DQNIMAGE_train.json', 'r') as f:
        data = json.load(f)

plt.figure(figsize=(16, 5))


plt.subplot(121)
plot_live_data(data['steps_done'], data['rewards'], avg=False,
               title={'title': "Rewards per Step", 'raw': 'Raw Rewards',
                       'smooth': 'Smoothed Rewards',
                       'acc' : 'Cumulative Average Rewards', 'xlab': 'Episode',
                       'ylab': 'Rewards'})

plt.subplot(122)
plot_train_data(data['steps_done'], data['td_error_sq'], avg=False,
                title={'title' : "TD-Error**2 per Step", 'raw': 'Raw
                    TD-Error**2',
                    'smooth': 'Smoothed TD-Error**2', 'acc': 'Cumulative
                    Average TD-Error**2',
                    'xlab' : 'Episode', 'ylab': 'TD-Error**2'})

plt.tight_layout()
plt.savefig('report/figures/DQNRGBLearning_steps.pdf')
plt.savefig('report/figures/DQNRGBLearning_steps.png')

```

```
plt.show()
```



```
[21]: rewards = data['rewards']
steps_taken = data['steps_taken']
finish_counter = data['finish_counter']
episodes = len(rewards)
print('===== TRAIN SUMMARY =====')
print(f'Episodes      : {episodes}')
print(f'Completion rate: {finish_counter / episodes}")
print(f"Average Reward : {sum(rewards) / episodes:.3f}")
print(f"Average steps  : {sum(steps_taken) / episodes:.3f}")
```

```
===== TRAIN SUMMARY =====
```

```
Episodes      : 1500
Completion rate: 0.98
Average Reward : 0.892
Average steps  : 30.068
```

0.6 Evaluation

```
[22]: from DQN.DQNIMAGE import eval as dqn_image_eval
from DQN.DQNIMAGE import CNN_DQN, FrameStack
from Utility.DQNEpsilonStrategy import EpsilonGreedy as DQNEpsilonGreedy
from Utility.MiniGrid import MiniGridImage, get_device
import torch

device = get_device()

env_dqn_img = MiniGridImage(render_mode='rgb_array')
policy_net_img = CNN_DQN(56, 56, 3).to(device)
```

```
[23]: summary(policy_net_img, input_size=(1,4,56,56), device=device,
    ↪ col_names=['input_size', 'kernel_size', 'output_size', 'num_params',
    ↪ 'trainable'])
```

```

[23]: =====
=====
=====
Layer (type:depth-idx)      Input Shape      Kernel Shape
Output Shape               Param #          Trainable
=====
=====
=====
CNN_DQN                    [1, 4, 56, 56]   --
[1, 3]                     --              True
  Sequential: 1-1          [1, 4, 56, 56]   --
[1, 512]                   --              True
  Conv2d: 2-1              [1, 4, 56, 56]   [3, 3]
[1, 16, 27, 27]           576             True
  BatchNorm2d: 2-2        [1, 16, 27, 27]   --
[1, 16, 27, 27]           32              True
  ReLU: 2-3               [1, 16, 27, 27]   --
[1, 16, 27, 27]           --              --
  Conv2d: 2-4              [1, 16, 27, 27]   [3, 3]
[1, 32, 13, 13]           4,608           True
  BatchNorm2d: 2-5        [1, 32, 13, 13]   --
[1, 32, 13, 13]           64              True
  ReLU: 2-6               [1, 32, 13, 13]   --
[1, 32, 13, 13]           --              --
  Conv2d: 2-7              [1, 32, 13, 13]   [3, 3]
[1, 64, 6, 6]             18,432          True
  BatchNorm2d: 2-8        [1, 64, 6, 6]     --
[1, 64, 6, 6]             128             True
  ReLU: 2-9               [1, 64, 6, 6]     --
[1, 64, 6, 6]             --              --
  Conv2d: 2-10             [1, 64, 6, 6]     [3, 3]
[1, 128, 2, 2]            73,728          True
  BatchNorm2d: 2-11       [1, 128, 2, 2]    --
[1, 128, 2, 2]            256             True
  Flatten: 2-12           [1, 128, 2, 2]    --
[1, 512]                   --              --
  Sequential: 1-2         [1, 512]          --
[1, 3]                     --              True
  Linear: 2-13            [1, 512]          --
[1, 64]                    32,832          True
  ReLU: 2-14              [1, 64]           --
[1, 64]                    --              --
  Linear: 2-15            [1, 64]           --
[1, 3]                     195             True
=====
=====
=====
=====

```

```
Total params: 130,851
Trainable params: 130,851
Non-trainable params: 0
Total mult-adds (M): 2.19
```

```
=====
=====
=====
Input size (MB): 0.05
Forward/backward pass size (MB): 0.32
Params size (MB): 0.52
Estimated Total Size (MB): 0.89
=====
=====
=====
```

```
[24]: policy_net_img.load_state_dict(torch.load('DQN/dqn_image.pth',
        ↪map_location=device))
      dqn_image_eval(env_dqn_img, policy_net_img, DQNEpsilonGreedy(0, 3), 1000)
```

Evaluation...

```
/Users/lucien/.pyenv/versions/3.10.14/envs/torch/lib/python3.10/site-
packages/gymnasium/core.py:311: UserWarning: WARN: env.get_frame to get
variables from other wrappers is deprecated and will be removed in v1.0, to get
this variable you can do `env.unwrapped.get_frame` for environment variables or
`env.get_wrapper_attr('get_frame')` that will search the reminding wrappers.
  logger.warn(
```

===== EVALUATION SUMMARY =====

```
Evaluation episodes: 1000
Completion rate      : 1.0
Average Reward      : 0.961
Average steps       : 11.000
```

```
[25]: imgs = []
      fig = plt.figure(figsize=(5, 5))
      frame_stack = FrameStack(56, 56, 4)
      strategy_dqn_img = DQNEpsilonGreedy(0, 3)
      current_state = env_dqn_img.reset()
      current_stack = frame_stack.push(current_state, True)
      plt.axis('off')
      imgs.append([plt.imshow(env_dqn_img.render(), animated=True)])
      for step in range(env_dqn_img.maxSteps):
          action = strategy_dqn_img.select_action(current_stack, policy_net_img)
          next_state, reward, done, truncated = env_dqn_img.step(action.item())
          next_stack = frame_stack.push(next_state, False)
          imgs.append([plt.imshow(env_dqn_img.render(), animated=True)])
```

```
    if done or truncated:
        break
    current_stack = next_stack
ani = animation.ArtistAnimation(fig, imgs, interval=100, repeat_delay=1000)
ani.save('DQN/dqn_img.gif', writer='imagemagick')
plt.close()
HTML(ani.to_jshtml())
```

MovieWriter imagemagick unavailable; using Pillow instead.

[25]: <IPython.core.display.HTML object>